

SHAQUONNA WILLIAMS

Software Engineer — ML Systems & Infrastructure

(443) 599-2756 · shaquonna.lynna@gmail.com · Baltimore, MD · [Portfolio](#) · [LinkedIn](#) · [GitHub](#)

SUMMARY

Full-stack software engineer with firsthand experience operating GPU-accelerated ML infrastructure, recovering broken CUDA/driver stacks, tuning a vLLM inference server for LLM serving, and running production API systems on GCP.

SKILLS

Languages: Python, JavaScript, TypeScript, SQL, HTML, CSS

ML & Systems Infrastructure: GCP (Compute Engine, GPU provisioning), Docker & Docker Compose, NVIDIA GPU instances (L4), CUDA/DKMS troubleshooting, vLLM (tensor parallelism, KV-cache tuning), Nginx, distributed systems, PostgreSQL, SQLite, SQLAlchemy, Alembic

Frameworks & APIs: FastAPI, Flask, Express.js, REST, React, Redux, NestJS, Prisma

Tooling & Practices: Bash, Git/GitHub, Kanban, Postman, Chrome DevTools, Mocha, Figma, Render

EXPERIENCE

Full-Stack Software Engineer

May 2025 – June 2026

Remote

- Architected a production-grade, GPU-accelerated staging environment on GCP (NVIDIA L4 instances) for an AI service, implementing automated scheduling policies that cut infrastructure costs by 75% while preserving development team availability during business hours.
- Diagnosed and recovered a broken NVIDIA driver/DKMS/CUDA stack after a GPU VM migration, rebuilding kernel modules against the running kernel and reconfiguring Nvidia-container-toolkit to restore Docker GPU passthrough for a self-hosted vLLM inference server, verified at 98–99% GPU utilization under load.
- Tuned vLLM serving parameters for a Qwen2.5-7B-Instruct deployment, scaling from 1 to 2 NVIDIA L4 GPUs with tensor parallelism, tripling max context length (9K → 32K tokens), and expanding concurrent request capacity roughly 10x through KV-cache, batching, and prefix-caching tuning.
- Deployed a containerized microservices application on GCP with Docker Compose orchestration, configuring a Nginx reverse proxy and SSL/TLS encryption to route traffic securely across a React frontend, FastAPI backend, and AI service endpoints.
- Built a secure API gateway with configurable rate limiting (10–100 req/s), multi-service request routing, and security headers (HSTS, CSP, X-Frame-Options) to harden the production environment against common web vulnerabilities.

Operator / Manager

Apr 2021 – Mar 2024

- Led strategic planning and operations for a growing small business, managing budgeting and financial reporting and driving a 25% reduction in overhead costs.

Independent Contract Driver

Apr 2015 – Present

- Applied route optimization and scheduling tools to manage logistics, improve on-time delivery, and reduce customer disputes.

PROJECTS

Eater-Verse

- Built a RESTful API and responsive Figma-based frontend as part of a 4-engineer team, using GitHub and Kanban to coordinate features and minimize merge conflicts.

Flexitude

- Designed a normalized SQLite database with many-to-many associations and fetched data directly into frontend components to keep the client stateless.

EDUCATION

App Academy — Software Engineer Certification, Full-Stack Software Engineering Bootcamp

2024 – 2025

Stratford University — A.A.S. in Advanced Culinary Arts

2013 – 2016

LANGUAGES

English — Fluent · Japanese — Working Knowledge